

```
In [1]: import torch
import torch.nn as nn
import torchvision
from torchvision import datasets, transforms, models
import numpy as np
import matplotlib.pyplot as plt
import random

device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")
```

```
In [33]: #Exercise 1
transforms=transforms.Compose([transforms.ToTensor(),\
                               transforms.Normalize((0.4376821, 0.4437697, 0.47280442),\
                                                     (0.19893012, 0.20101562, 0.19703614))])

train_set = torchvision.datasets.SVHN(root='./data',download=True, transform=transforms)
train_size = len(train_set)//10
validation_size = len(train_set) - train_size
train_set, validation_set=torch.utils.data.random_split(train_set,[train_size, validation_size])
dataset_sizes={'train':train_size,'val':validation_size}
test_set = torchvision.datasets.SVHN(root='./data',split='test',download=True, transform=transforms)
dataset_sizes['train']=train_set,'val':validation_set,'test':test_set]
dataloaders = {x: torch.utils.data.DataLoader(datasets[x], batch_size=128, shuffle=True, num_workers=2)
               for x in ['train'], 'val'],'test'}}

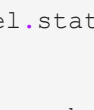
Using downloaded and verified file: ./data/train_32x32.mat
Using downloaded and verified file: ./data/test_32x32.mat
```

```
In [ ]: def imshow(inp, title=None):
    """Imshow for Tensor."""
    inp = inp.numpy().transpose((1, 2, 0))
    plt.imshow(inp)
    if title is not None:
        plt.title(title)
    plt.pause(0.001) # pause a bit so that plots are updated

# Get a batch of training data
inputs, classes = next(iter(dataloaders['train']))

# Make a grid from batch
out = torchvision.utils.make_grid(inputs)

imshow(out, title=[x for x in classes])
```



```
In [3]: #Exercise 2
input_size=32*32*3
output_size=10
def Fcn(l1_size=128,l2_size=64):
    return nn.Sequential(nn.Linear(input_size, l1_size),
                        nn.Linear(l1_size, l2_size),
                        nn.ReLU(),
                        nn.Linear(l2_size, output_size),
                        nn.ReLU())

def Cnn(ker_size=3):
    con_out_size=((32-ker_size+3)//2)
    con_out_size=(con_out_size-ker_size+3)//2)
    con_out_size=con_out_size**2
    return nn.Sequential(nn.Conv2d(in_channels=3,out_channels=10,
                                   kernel_size=ker_size,padding=1,stride=1),
                        nn.ReLU(),
                        nn.MaxPool2d(kernel_size=2,stride=2,padding=0),
                        nn.Conv2d(in_channels=10,out_channels=20,
                                   kernel_size=ker_size,padding=1,stride=1),
                        nn.ReLU(),
                        nn.MaxPool2d(kernel_size=2,stride=2,padding=0),
                        nn.Flatten(),
                        nn.Linear(con_out_size*20, 64),
                        nn.ReLU(),
                        nn.Linear(64,output_size))
```

```
In [10]: #Exercise 3/4
import copy
def train_model(model, criterion, optimizer,fc=True,num_epochs=30,debug=0):
    best_model_wts = copy.deepcopy(model.state_dict())
    best_acc = 0.0
    model=model.to(device)
    epoch_loss = {'train':np.zeros(num_epochs),'val':np.zeros(num_epochs)}
    epoch_acc = {'train':np.zeros(num_epochs),'val':np.zeros(num_epochs)}
    for epoch in range(num_epochs):
        if debug==1 or epoch == num_epochs-1:
            print('Epoch {}/{}'.format(epoch, num_epochs - 1))
            print('-' * 10)

        # Each epoch has a training and validation phase
        for phase in ['train','val']:
            if phase == 'train':
                model.train() # Set model to training mode
            else:
                model.eval() # Set model to evaluate mode

            running_loss = 0.0
            running_corrects = 0

            # Iterate over data.
            for inputs, labels in dataloaders[phase]:
                if fc:
                    inputs = inputs.view(inputs.shape[0], -1)
                    labels = labels.to(device)
                    # zero the parameter gradients
                    optimizer.zero_grad()

                    # forward
                    # track history if only in train
                    with torch.set_grad_enabled(phase == 'train'):
                        outputs = model(inputs)
                        preds = torch.argmax(outputs,dim=1)
                        loss = criterion(outputs, labels)

                    # backward + optimize only if in training phase
                    if phase == 'train':
                        loss.backward()
                        optimizer.step()

            # statistics
            running_loss += loss.item() * inputs.size(0)
            running_corrects += torch.sum(preds == labels.data)

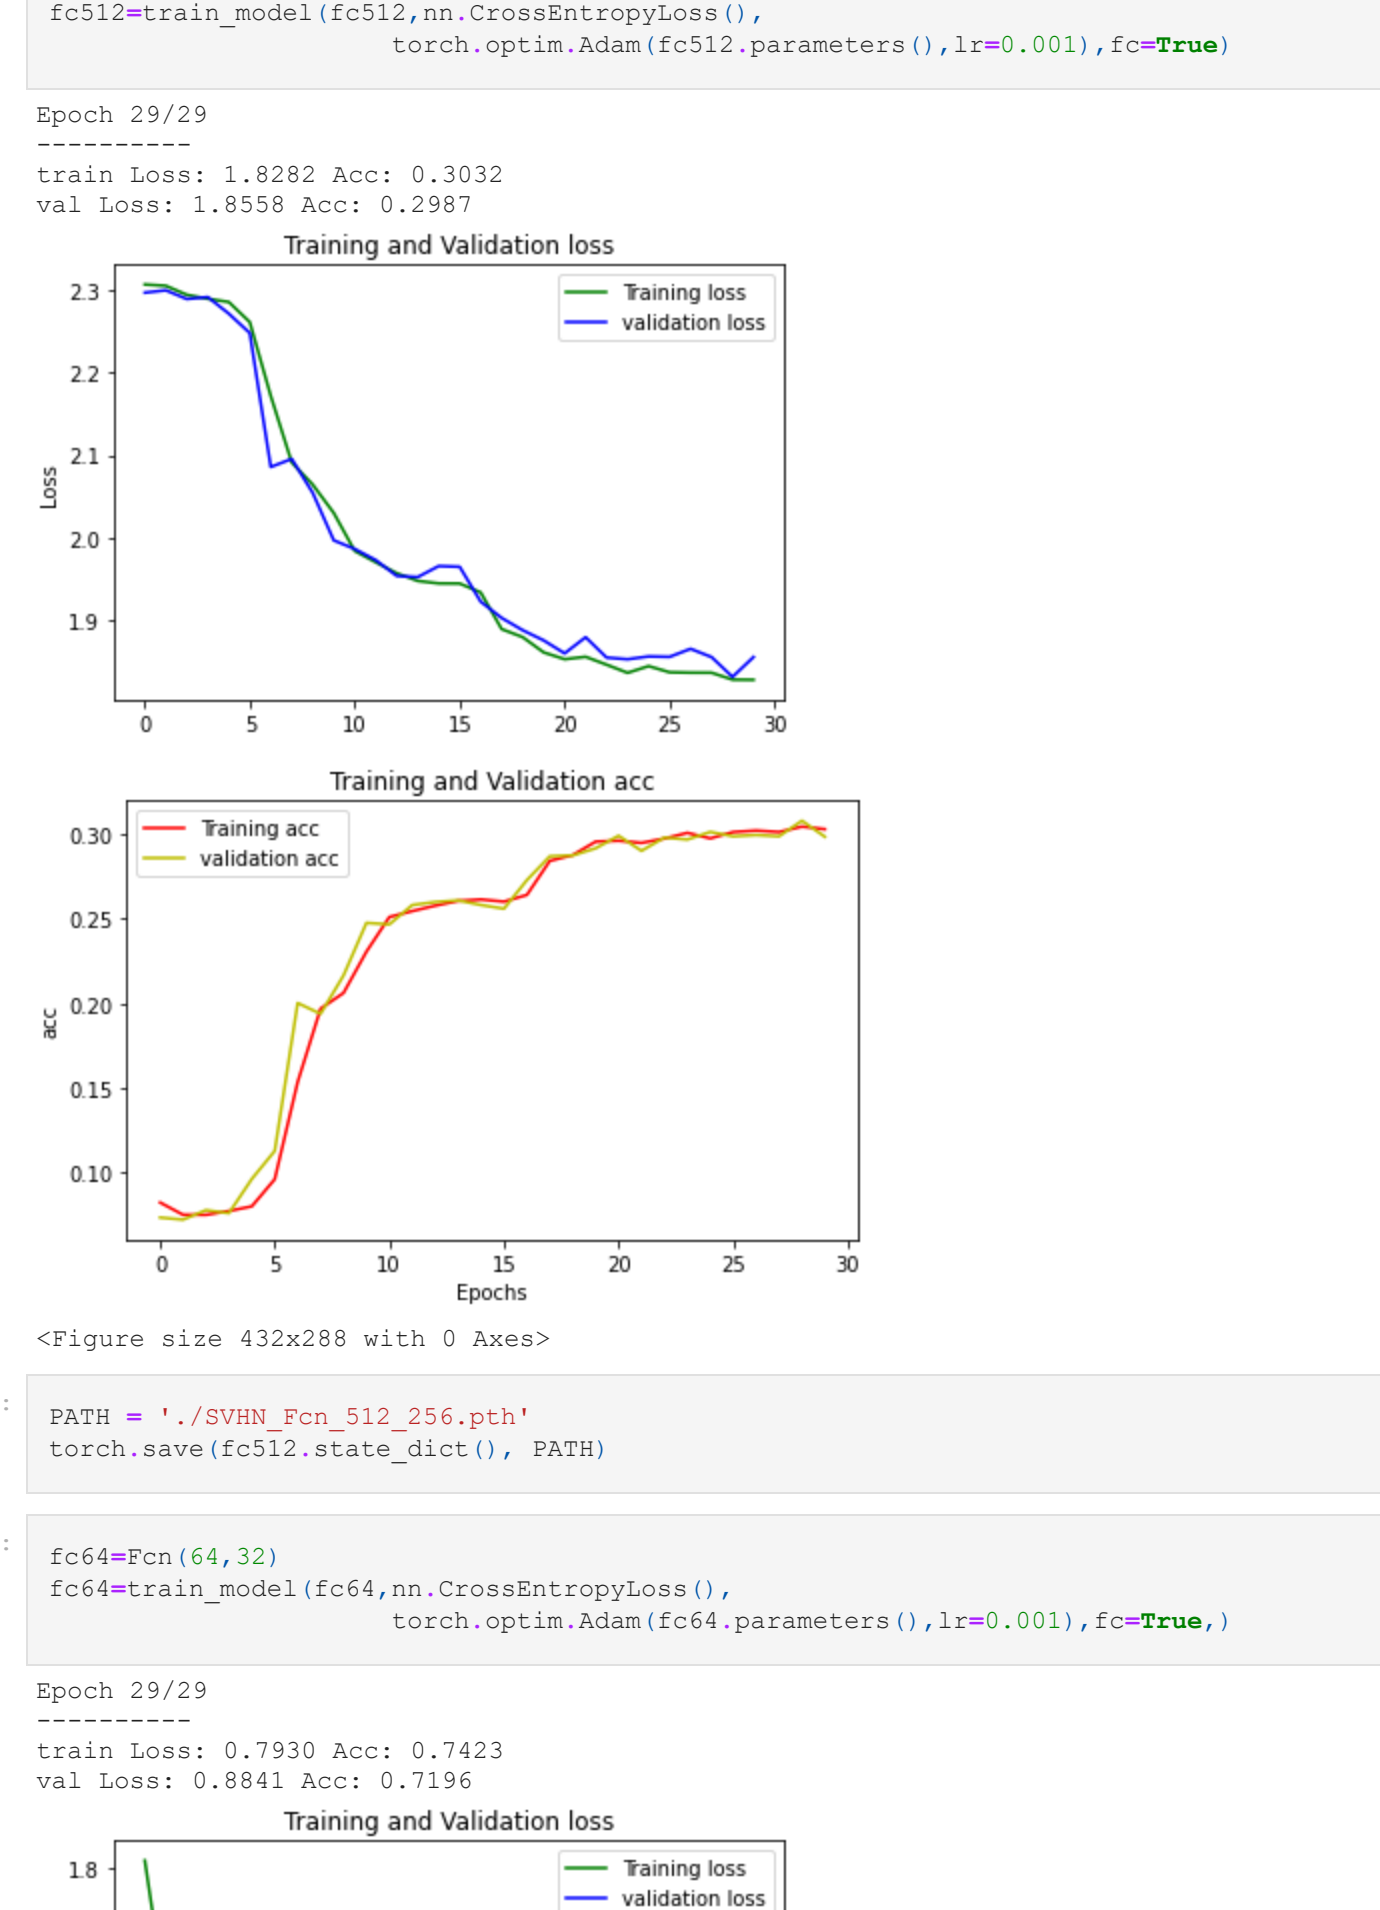
            epoch_loss[phase][epoch] = running_loss / dataset_sizes[phase]
            epoch_acc[phase][epoch] = running_corrects.double() / dataset_sizes[phase]

        if phase == 'train' and (debug==1 or epoch == num_epochs-1):
            print('{} Loss: {:.4f} Acc: {:.4f}'.format(
                phase, epoch_loss[phase][epoch], epoch_acc[phase][epoch]))
        elif phase == 'val' and (debug==1 or epoch == num_epochs-1):
            print('{} Loss: {:.4f} Acc: {:.4f}'.format(
                phase, epoch_loss[phase][epoch], epoch_acc[phase][epoch]))

    plt.plot(range(num_epochs), epoch_loss['train'], 'g', label='Training loss')
    plt.plot(range(num_epochs), epoch_loss['val'], 'b', label='validation loss')
    plt.title('Training and Validation loss')
    plt.ylabel('Loss')
    plt.legend()
    plt.show()
    plt.clf()
    plt.plot(range(num_epochs), epoch_acc['train'], 'r', label='Training acc')
    plt.plot(range(num_epochs), epoch_acc['val'], 'y', label='validation acc')
    plt.title('Training and Validation acc')
    plt.xlabel('Epochs')
    plt.ylabel('acc')
    plt.legend()
    plt.tight_layout()
    plt.show()
    plt.clf()
    return model
```

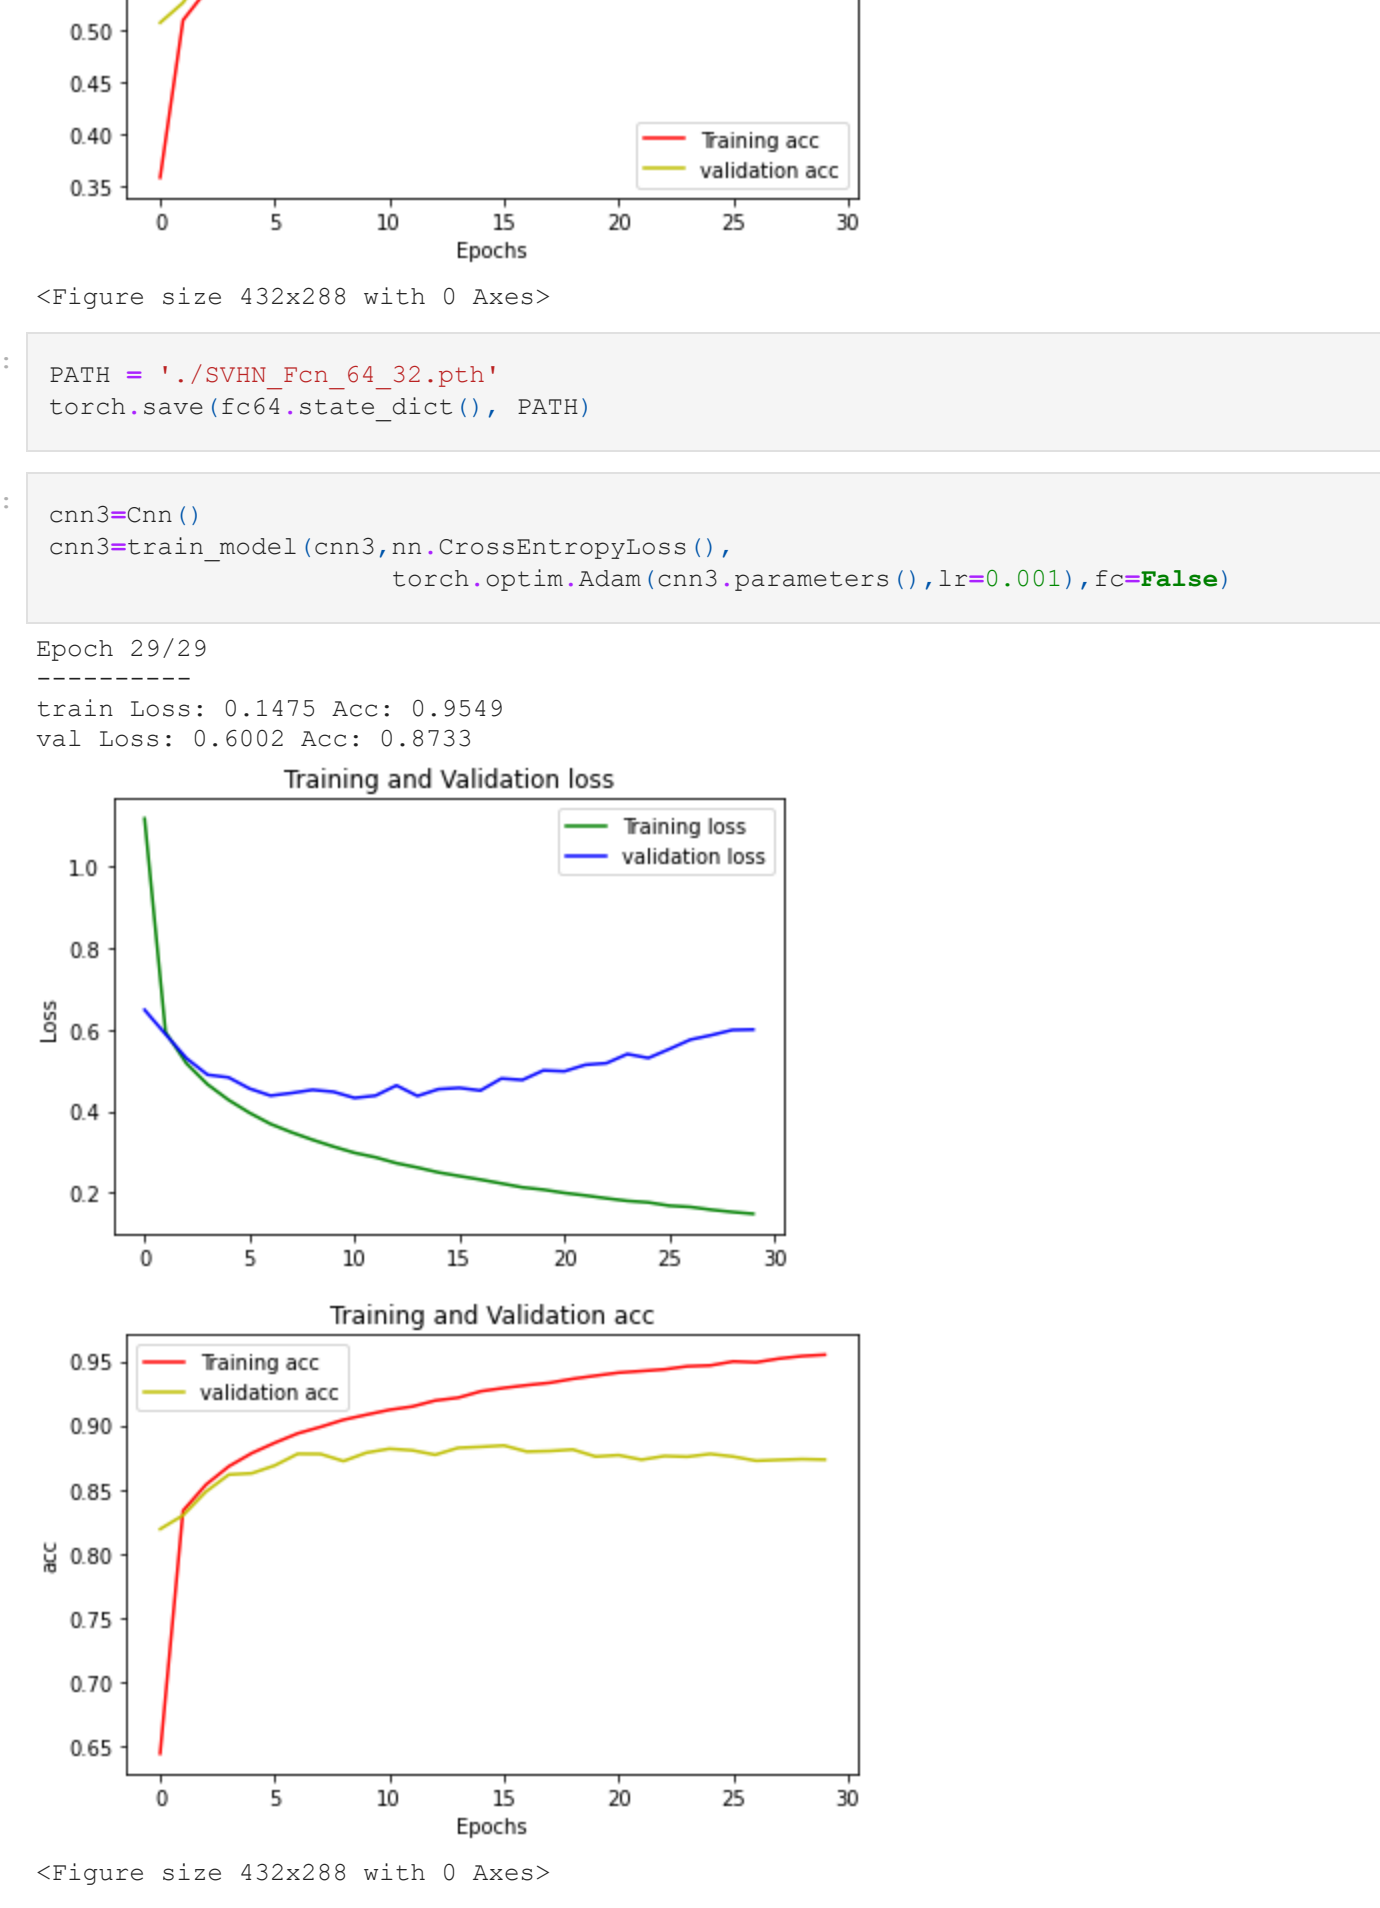
Exercises 3+4

```
In [11]: fc128=Fcn()
fc128=train_model(fc128,nn.CrossEntropyLoss(),
                  torch.optim.Adam(fc128.parameters(),lr=0.001),fc=True)
```



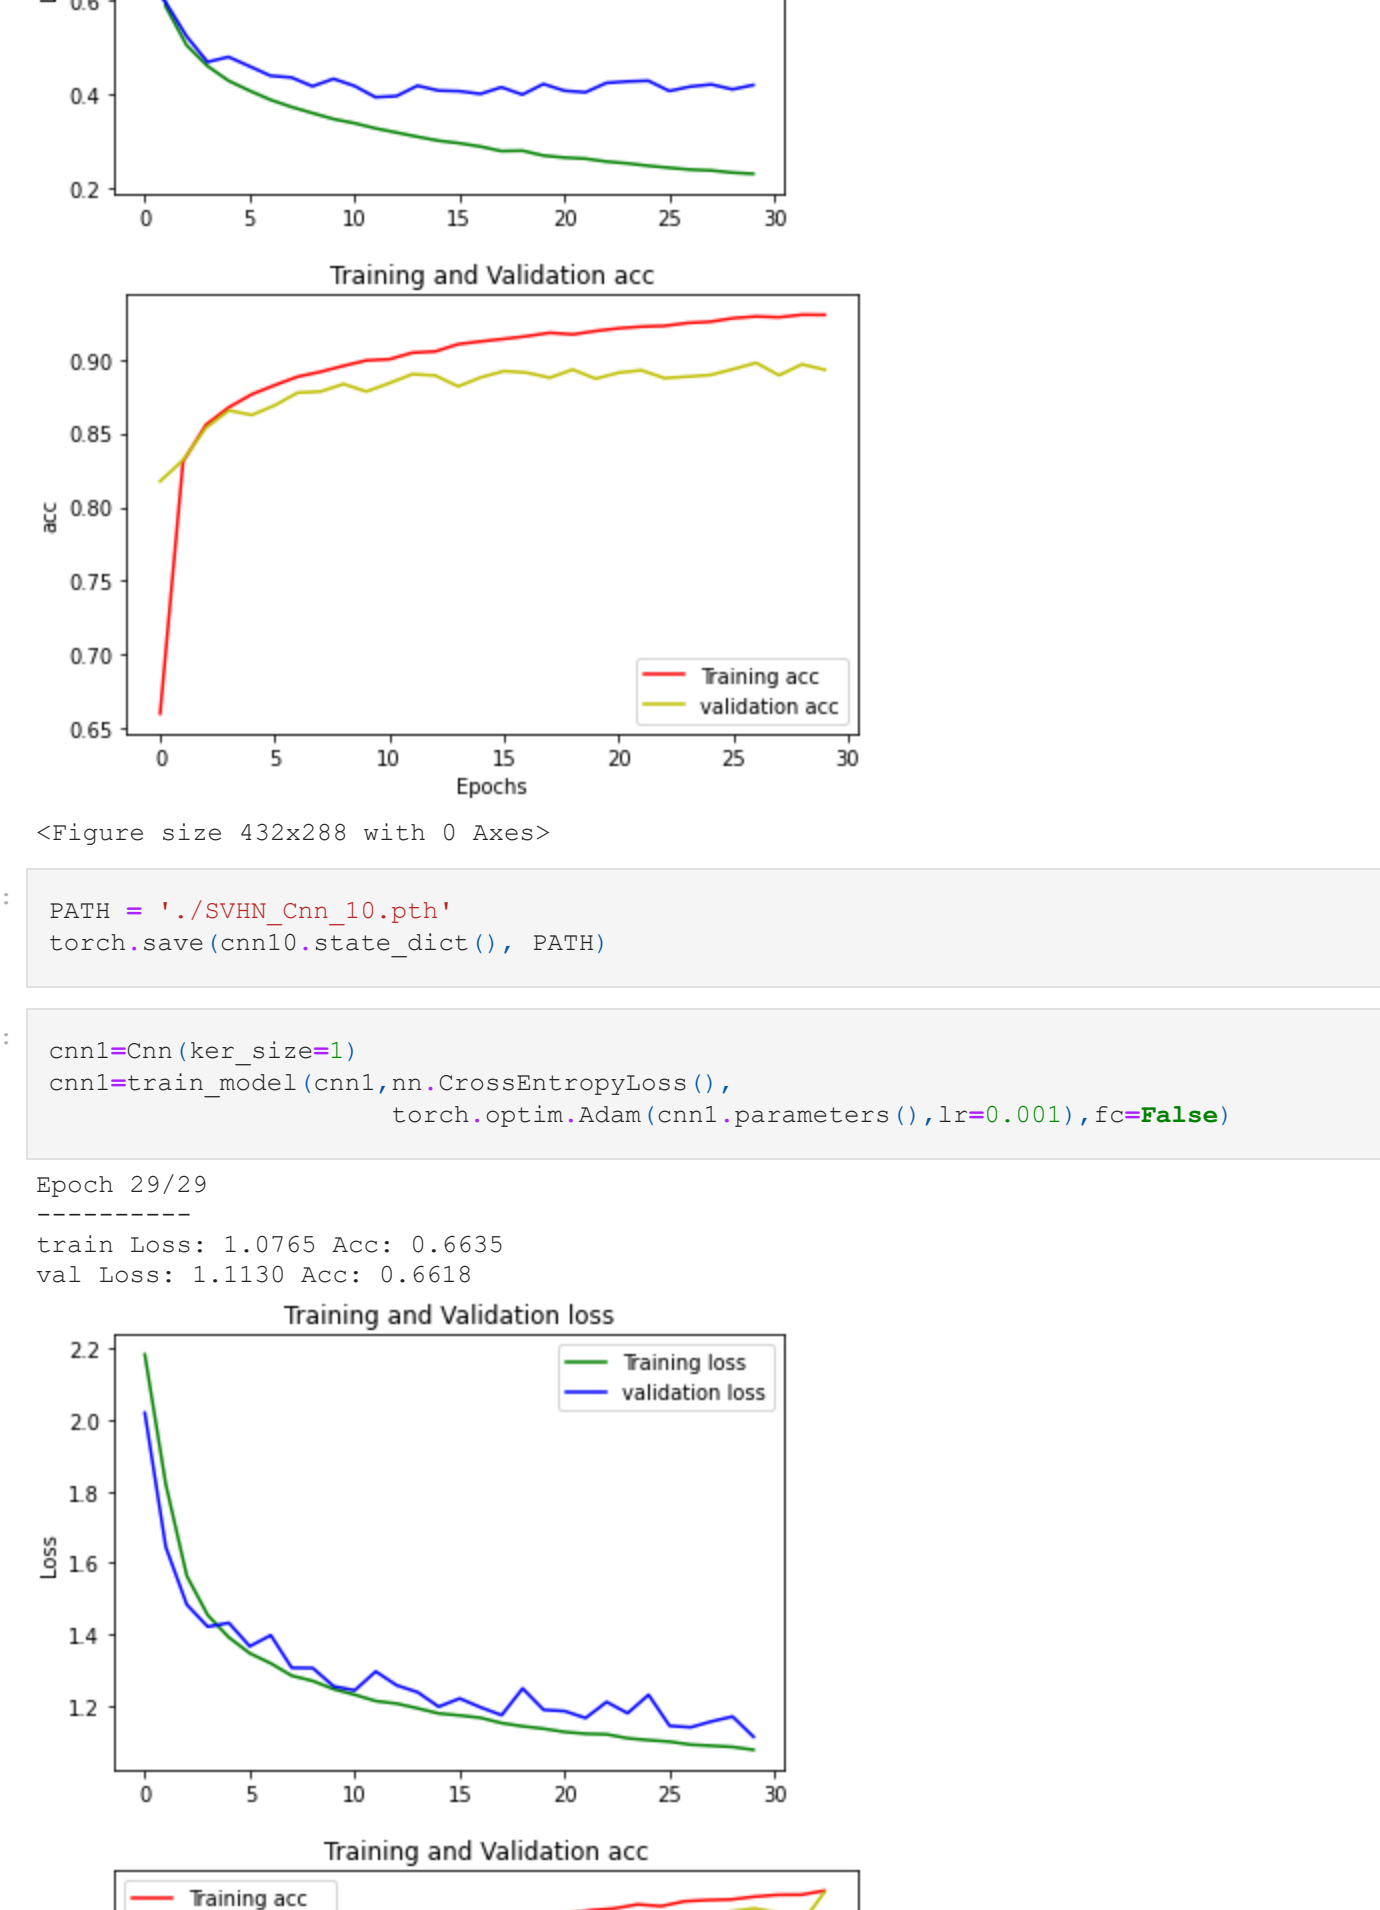
```
In [12]: PATH = './SVHN_Fcn_128_64.pth'
torch.save(fc128.state_dict(), PATH)
```

```
In [13]: fc512=Fcn(512,256)
fc512=train_model(fc512,nn.CrossEntropyLoss(),
                  torch.optim.Adam(fc512.parameters(),lr=0.001),fc=True)
```



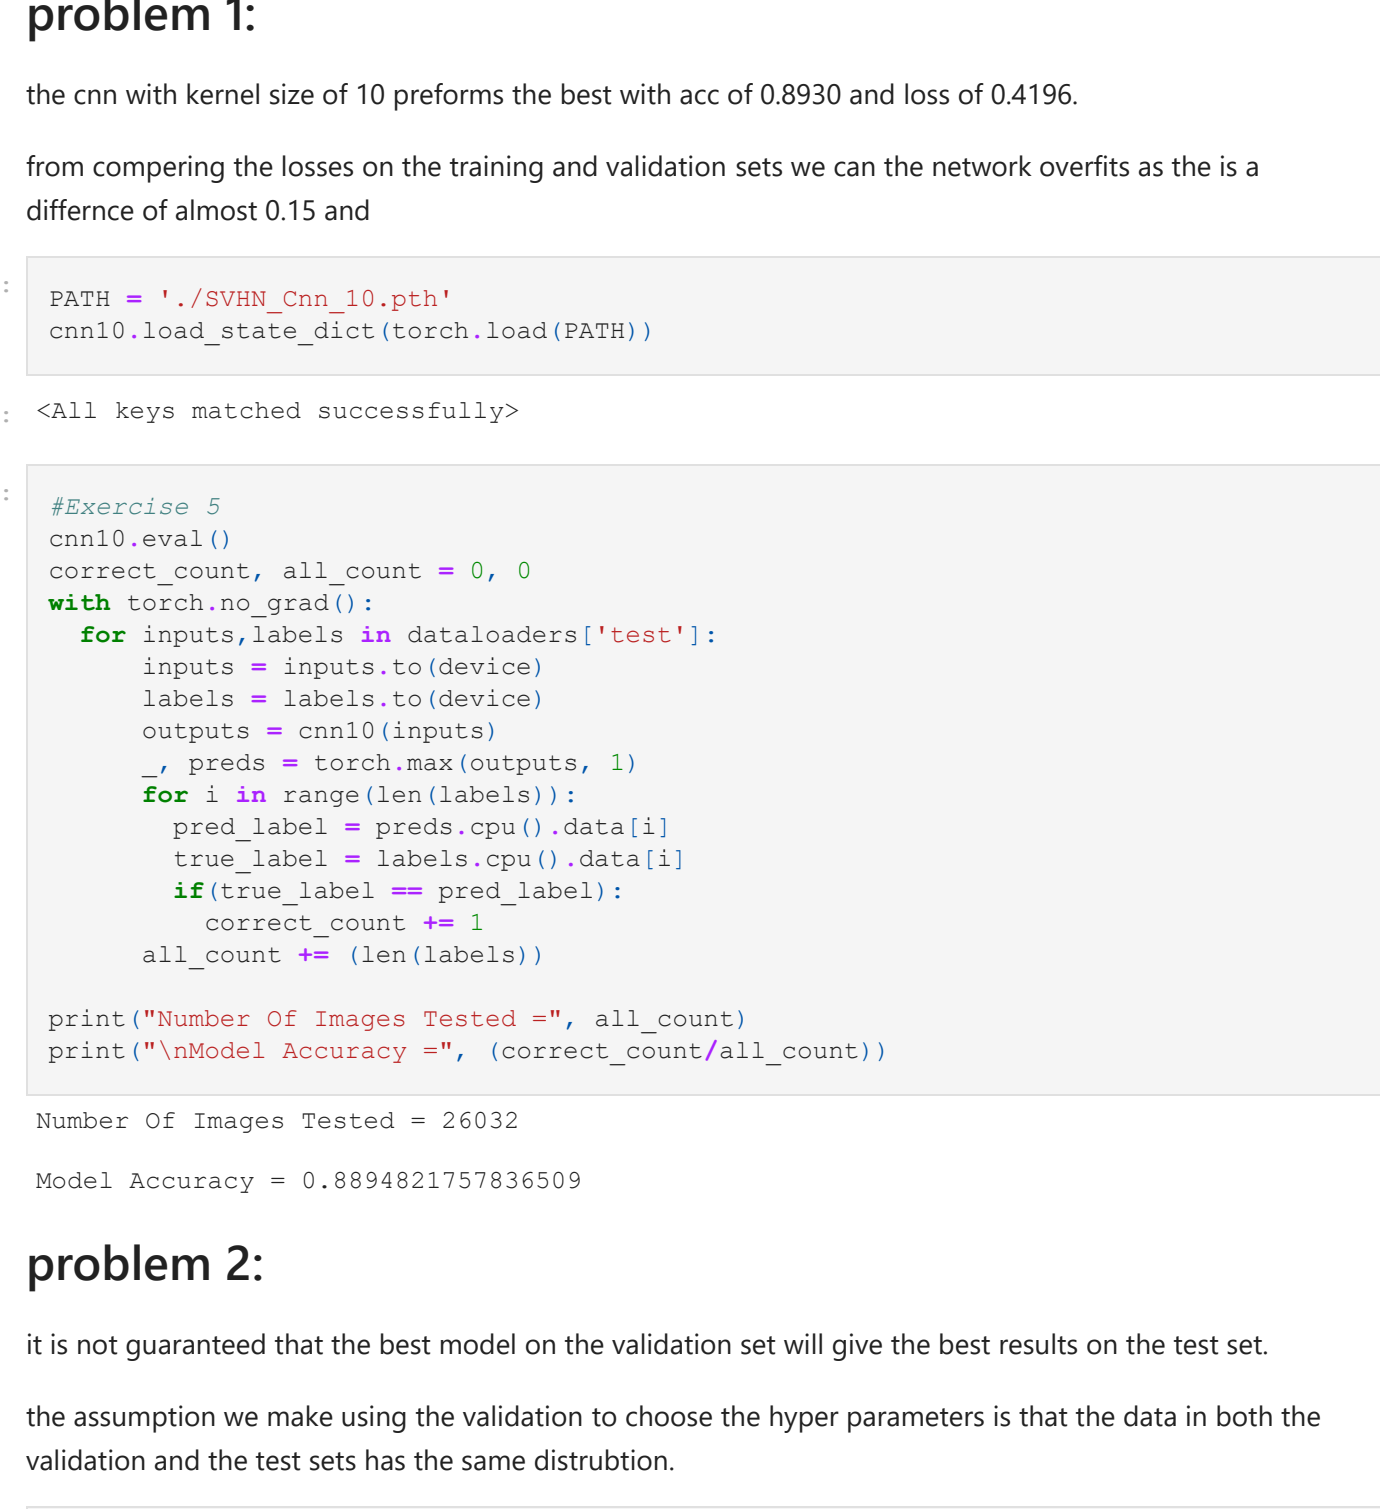
```
In [14]: PATH = './SVHN_Fcn_512_256.pth'
torch.save(fc512.state_dict(), PATH)
```

```
In [27]: fc64=Fcn(64,32)
fc64=train_model(fc64,nn.CrossEntropyLoss(),
                  torch.optim.Adam(fc64.parameters(),lr=0.001),fc=True,)
```



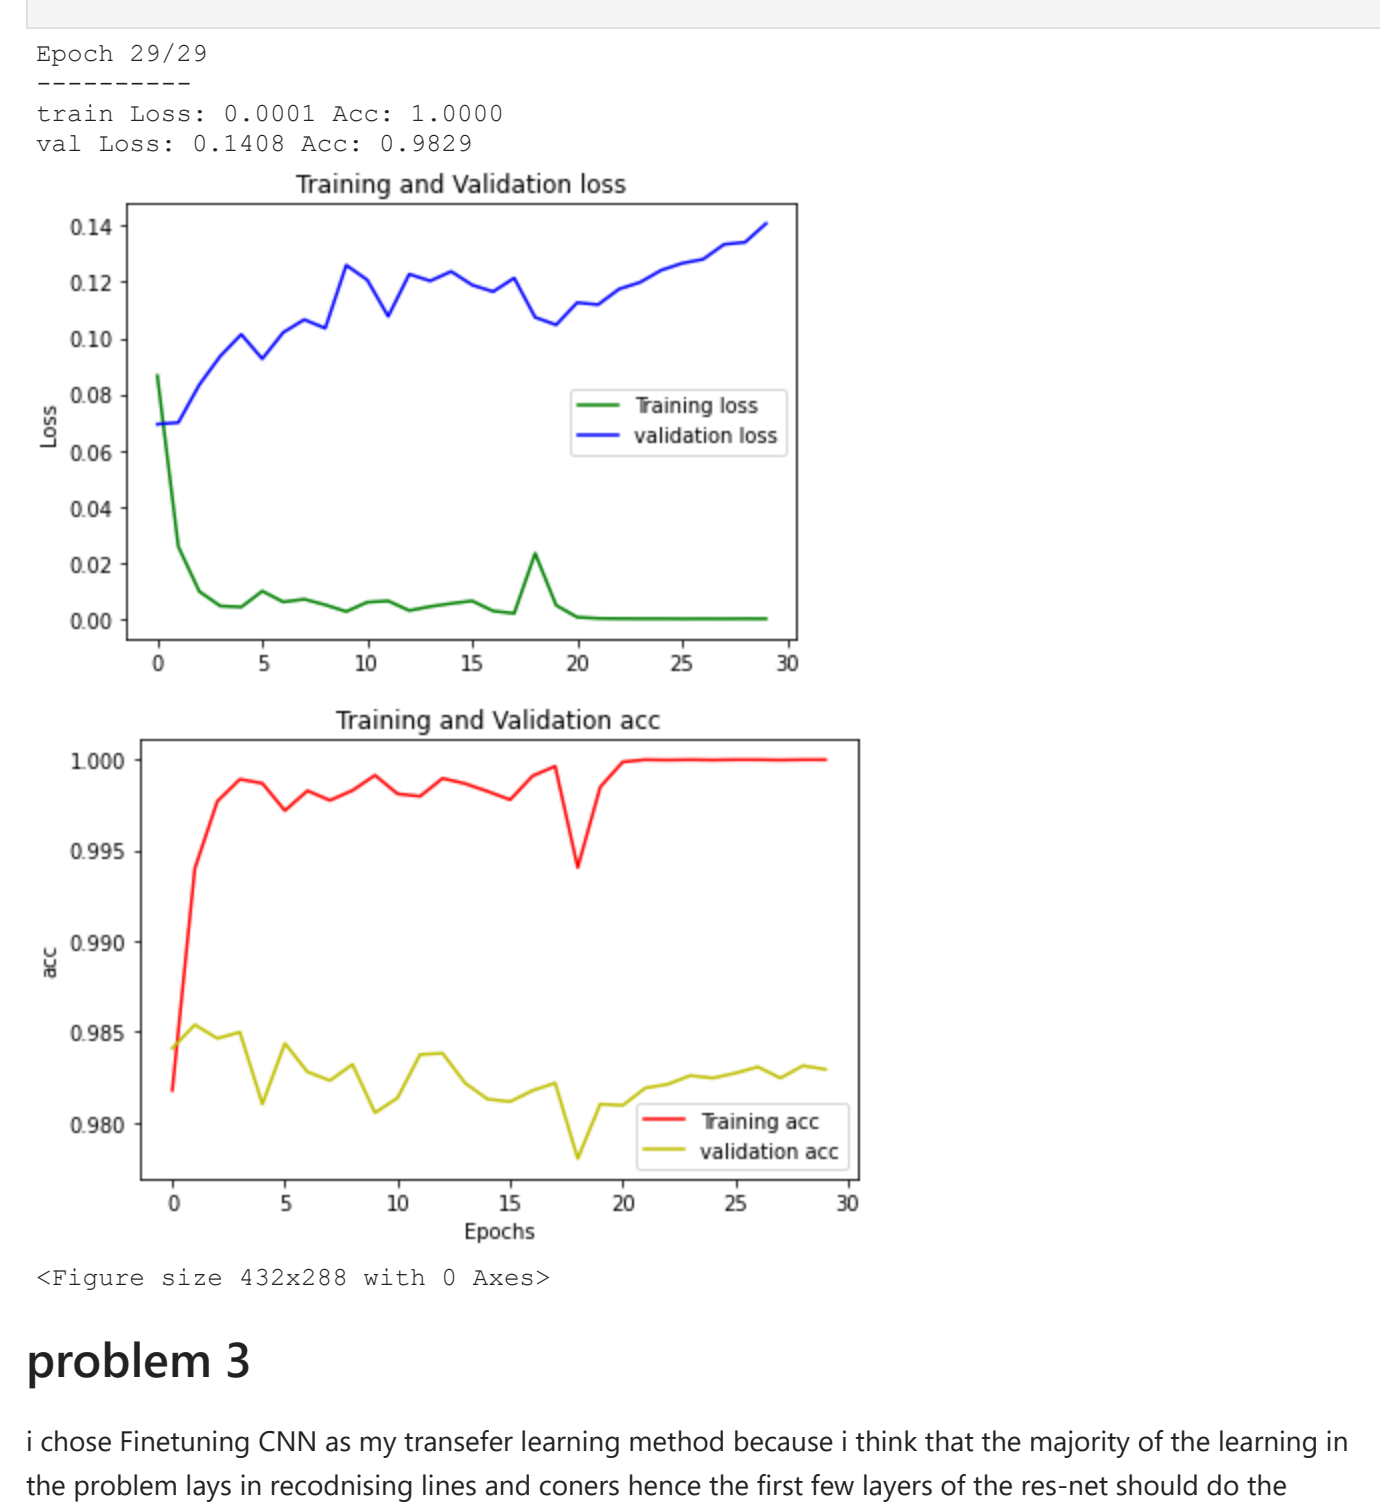
```
In [28]: PATH = './SVHN_Fcn_64_32.pth'
torch.save(fc64.state_dict(), PATH)
```

```
In [17]: cnn3=Cnn()
cnn3=train_model(cnn3,nn.CrossEntropyLoss(),
                  torch.optim.Adam(cnn3.parameters(),lr=0.001),fc=False)
```



```
In [18]: PATH = './SVHN_Cnn_3.pth'
torch.save(cnn3.state_dict(), PATH)
```

```
In [19]: cnn10=Cnn(ker_size=10)
cnn10=train_model(cnn10,nn.CrossEntropyLoss(),
                  torch.optim.Adam(cnn10.parameters(),lr=0.001),fc=False)
```



```
In [20]: PATH = './SVHN_Cnn_10.pth'
torch.save(cnn10.state_dict(), PATH)
```

```
In [21]: cnn1=Cnn(ker_size=1)
cnn1=train_model(cnn1,nn.CrossEntropyLoss(),
                  torch.optim.Adam(cnn1.parameters(),lr=0.001),fc=False)
```



```
In [22]: PATH = './SVHN_Cnn_1.pth'
torch.save(cnn1.state_dict(), PATH)
```

problem 1:

the cnn with kernel size of 10 preforms the best with acc of 0.8930 and loss of 0.4196.

from compering the losses on the training and validation sets we can the network overfits as the is a difference of almost 0.15 and

```
In [23]: PATH = './SVHN_Cnn_10.pth'
cnn10.load_state_dict(torch.load(PATH))

Out[23]: <All keys matched successfully>
```

```
In [24]: #Exercise 5
cnn10.eval()
correct_count, all_count = 0, 0
with torch.no_grad():
    for inputs,labels in dataloaders['test']:
        inputs = inputs.to(device)
        labels = labels.to(device)
        outputs = cnn10(inputs)
        preds = torch.max(outputs, 1)
        for i in range(len(labels)):
            pred_label = preds.cpu().data[i]
            true_label = labels.cpu().data[i]
            if true_label == pred_label:
                correct_count += 1
        all_count += (len(labels))

print("Number Of Images Tested = ", all_count)
print("\nModel Accuracy = ", (correct_count/all_count))
```

Number Of Images Tested = 26032
Model Accuracy = 0.8894821757836509

problem 2:

it is not guaranteed that the best model on the validation set will give the best results on the test set.

the assumption we make using the validation to choose the hyper parameters is that the data in both the validation and the test sets has the same distribution.

```
In [25]: #Exercise 6
model_conv = torchvision.models.resnet18(pretrained=True)

# Parameters of newly constructed modules have requires_grad=True by default
num_fters = model_conv.fc.in_features
model_conv.fc = nn.Linear(num_fters, 10)

model_conv = model_conv.to(device)

criterion = nn.CrossEntropyLoss()

# Observe that only parameters of final layer are being optimized as
# opposed to before.
optimizer_conv = torch.optim.Adam(model_conv.parameters(),lr=0.001)
```

download: <https://download.pytorch.org/models/resnet18-5c106cd6.pth> to /root/.cache/torch/hub/checkpoints/resnet18-5c106cd6.pth

```
In [34]: model_conv = train_model(model_conv, criterion, optimizer_conv,fc=False)
```


problem 3

i chose Finetuning CNN as my transfer learning method because i think that the majority of the learning in the problem lays in recognising lines and coners hence the first few layers of the res-net should do the majority of the work. Also there is a major difference between the images in imgnnet and our dataset because of that i believe that CNN as a fixed feature extractor wouldnt preform as good.

problem 4

the results of the new network are better then the results of all the CNN an FC networks i've implemented. those are the results on the validation set

1. best CNN - Loss: 0.4196 Acc: 0.8930
2. best FC - Loss: 0.6124 Acc: 0.8329
3. new net - Loss: 0.1408 Acc: 0.9829

i believe the depth of the network along side with residual block allow the res-net to learn better without vanishing gradient problem therefore it can solve the problem better

problem 5

```
In [ ]: #Exercise 6
for lr in [0.5,0.05,0.025,0.01,0.001,0.0001,0.000001]:
    print('lr= {}'.format(lr))
    print('-' * 10)
    model_conv = torchvision.models.resnet18(pretrained=True)

    # Parameters of newly constructed modules have requires_grad=True by default
    num_fters = model_conv.fc.in_features
    model_conv.fc = nn.Linear(num_fters, 10)

    model_conv = model_conv.to(device)

    criterion = nn.CrossEntropyLoss()

    optimizer_conv = torch.optim.Adam(model_conv.parameters(),lr=lr)
    epoch_loss = np.zeros(30)
    for epoch in range(30):
        model_conv.train() # Set model to training mode
        running_loss = 0.0
        running_corrects = 0

        # Iterate over data.
        for inputs, labels in dataloaders['train']:
            inputs = inputs.to(device)
            labels = labels.to(device)
            # zero the parameter gradients
            optimizer_conv.zero_grad()

            # forward
            # track history if only in train
            with torch.set_grad_enabled(True):
                outputs = model_conv(inputs)
                preds = torch.argmax(outputs,dim=1)
                loss = criterion(outputs, labels)
                loss.backward()
                optimizer_conv.step()

        # statistics
        running_loss += loss.item() * inputs.size(0)
        running_corrects += torch.sum(preds == labels.data)

        epoch_loss[epoch] = running_loss / dataset_sizes['train']

    plt.plot(range(30), epoch_loss, label='lr={}'.format(lr))

plt.title('Training loss')
plt.ylabel('Loss')
plt.legend()
plt.show()
```


As we can see in the graph, having a very high learning rate leads to the network not learning at all as the steps in the gradient direction are too big, so the network doesnt converge into any kind of minimum.

when the learning rate is to low the steps in the direction of the gradient are too small so there is a small improvement but the chances of converging into local minimum increase. for the rates in between the learning rate changes the speed of converages but the differences are that big meaning the steps are of a good size